

HEP_(e)xML

Intro to Deep Learning

Johannes Wagner

Last week recap

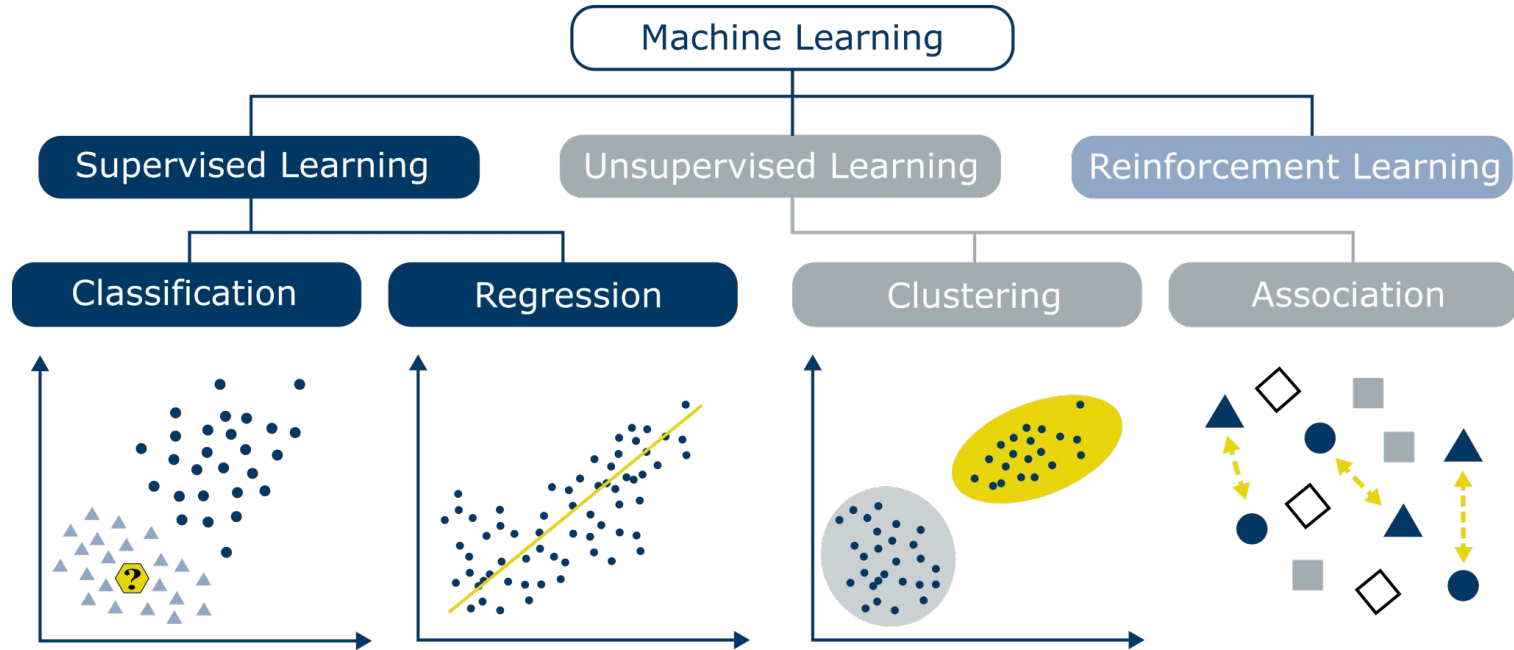
- We went through some of the basics of experimental particle physics
 - Particle detection, reconstruction and analysis
- Hands-on exercise showcased numpy/pyTorch fundamentals
- Today: Introduction of machine learning
 - Lecture will focus on broad concepts in ML and deep learning
 - Hands-on exercise will allow you to implement a simple deep neural network and apply it to a physics task

What is machine learning?

- Machine learning (ML) is a field that lives in the intersection between computer science and statistics
 - Related to, but not synonymous with artificial intelligence (AI)
- Problem: Given some data, how do we make predictions about them?
 - Need some function that maps our inputs to outputs that we are interested in
- Machine learning focuses on algorithms that can learn these functions without explicit programming

Types of ML problems

Source: [Towards Data Science](#)



Supervised learning


- Algorithm learns association of input data to known output labels
- Requires labeled training data
 - In particle physics, this comes from MC simulations
- Classification: place input data into discrete predefined categories
 - Commonly needed in image processing
- Regression: determine relationship between input data and continuous output
 - Linear regression is one of the simplest forms of ML

Unsupervised learning

- Algorithm learns to make associations between input data
 - Patterns are learned more broadly
- Does not require labeled data
- Clustering: group inputs into a set number of “clusters”
 - Like classification but clusters are not predefined
- Anomaly detection: spot elements that are anomalous compared to the typical element
 - Useful in BSM searches where you don't necessarily know what you're looking for

Reinforcement learning

- Requires an agent and some predefined task we want to accomplish
 - This task represents a “goal state”
- Our agent has a space of actions they can take to change their state
 - Need a function to define the distance between their current state and the goal state
- Often used in robotics and with games



Source: [KDNuggets](#)

The beginnings of ML

- ML has its beginnings in the 1950's with the advent of computers
 - Desire to build artificial brain lead to development of neural network algorithms
- Invention of the “Perceptron” in 1958 constitutes the first artificial neural network
 - Perceptron is a linear binary classification algorithm
- Development of many ML algorithms throughout the 20th century
 - Various different kinds of neural network algorithms - focus of this workshop
 - Other ML algorithms such as k-nearest neighbors, decision trees, etc. (see CS 189A)

How a Perceptron “learns”


- Given some input data $\{\mathbf{x}_j\}$, a perceptron predicts output class labels $\{y_j\}$ where $y_j(t) = f[\mathbf{w}(t) \cdot \mathbf{x}_j]$ and $f(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0, \\ 0 & \text{otherwise} \end{cases}$
- b is called the bias and shifts the decision boundary
- \mathbf{w} represents a trainable vector of weights, which are updated by the rule $w_i(t+1) = w_i(t) + r \cdot (d_j - y_j(t))x_{j,i}$, where d_j are the true labels
- This process is repeated until the \mathbf{w} vector no longer changes
- This is simple, but not very useful unless we have a problem that can be solved with a linear function

The “deep learning” revolution

- “Deep neural networks” are the most basic “deep learning” algorithms
 - Generalization of perceptrons that were not feasible to use until ~15 years ago due to lack of computing power
 - They are able to express arbitrary functions, not just linear ones
- Idea: stack multiple linear layers (i.e. perceptrons) together and connect them via “activation functions”
- Updating the weights to train the network is harder since the relationship between inputs and outputs is more complicated

Overview of DNN's


- Each node in the input layer represents one input feature
- Each arrow in the diagram represents a weighted connection to the next layer
- At each node, incoming weighted connections are summed and passed through an activation function to the next layer




Source: [Towards Data Science](#)

Activation functions

- Activation functions allow deep neural networks to model nonlinear functions
 - Successive linear matrix multiplications are still linear
- Should ideally be continuously differentiable
 - It's advantageous if the gradients don't tend to 0 or blow up too quickly anywhere in the domain
- Popular choices are ReLU or Sigmoid functions
 - Often used for regression/classification respectively



ReLU




Sigmoid

The forward pass

- Evaluation consists of matrix algebra
- Weights are updated based on backpropagation gradient descent

Source: [Towards Data Science](#)



Loss functions

- In order to understand how to update our network, we need some measure of how wrong our current predictions are: Loss functions!
 - Function that evaluates to a high number when our prediction is far off and a low number when it's close
- Classification: binary cross-entropy loss

$$L_{BCE} = -\frac{1}{N} [\sum_{j=1}^N y_j \log(\hat{y}_j) + (1 - y_j) \log(1 - \hat{y}_j)]$$

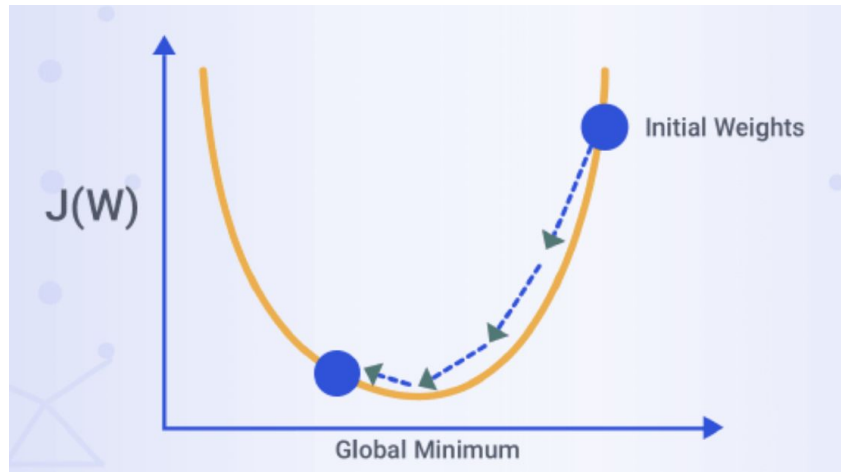
- Regression: mean-squared error loss

$$L_{MSE} = \frac{1}{N} \sum_{j=1}^N (\hat{y}_j - y_j)^2$$

Backpropagation gradient descent


- Training involves minimizing loss function
 - Loss is a function of NN weights, biases
- If we take the gradient with respect to the weights and biases, we get a set of directions we need to take a step in for loss minimization
 - Learning rate determines step size
- Usually this is done by taking the average gradient for some “batch” of data

Source: akira.ai



The loss landscape

- Loss landscape is a many-dimensional space (lots of weights in a NN) with a very complicated structure
 - Lots of local minima
- Need to find reliable ways to approach the global minimum without getting stuck
 - Choose the right learning rate



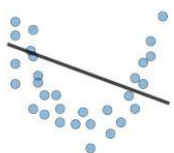


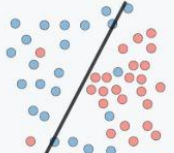
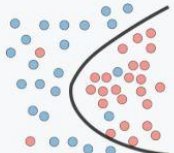
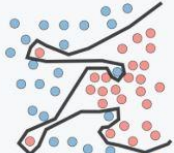


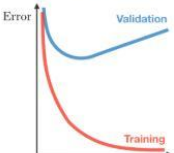
Source: [arxiv](#)

Beyond standard gradient descent

- Optimizers can use additional techniques to make loss minimization more effective
- Momentum: Incorporate previous gradient values into calculation of next step - keep “momentum” in some direction over time
- Adaptive gradients: Adapt step size along each dimension automatically based on smoothness of past gradients (smoother = larger step size)
- ADAM optimizer used commonly in modern applications
 - Gradient descent + momentum + adaptive gradients

Overfitting

- Overfitting happens when we train our neural network to match our training data too well
- Our data will have statistical fluctuations -> we don't want to fit these!
- Can check for this by separating out a slice of the dataset for “validation”
 - Other techniques such as “dropout” also help

| | Underfitting | Just right | Overfitting |
|-----------------------------|--|---|---|
| Symptoms | <ul style="list-style-type: none">• High training error• Training error close to test error• High bias | <ul style="list-style-type: none">• Training error slightly lower than test error | <ul style="list-style-type: none">• Very low training error• Training error much lower than test error• High variance |
| Regression illustration |  |  |  |
| Classification illustration |  |  |  |
| Deep learning illustration |  |  |  |
| Possible remedies | <ul style="list-style-type: none">• Complexify model• Add more features• Train longer | | <ul style="list-style-type: none">• Perform regularization• Get more data |

Source: [Kaggle](#)

Training data selection

- Once we have a trained neural network, we can apply it to new data we don't have the labels for!
- Neural network predictions are entirely dependent on what its shown during training
 - It won't be able to generalize if given an unreliable/biased training sample
- Carefully selecting your training data is one of the most important tasks when developing a NN algorithm

Recap: What is a deep neural network?

- A highly nonlinear function with many parameters that takes a vector as input and returns a scalar or vector
- A loose model of the neuronal connections in a human brain
- A trainable algorithm that can be applied to data to extract some relevant quantity from its inputs

Hands-on exercise #2

Deep Neural Networks

Exercise agenda

- Exercise #1: [pyTorch](#) (Credit: pyTorch Team)
 - We will finish the rest of the chapters starting from “Build Model”
- Exercise #2: “[Intro to Deep Learning](#)” Jupyter notebook
- Please take a moment to fill out the [feedback form](#) when you’re done